
CS589: Machine Learning - Fall 2017

Homework 2: Classification

Assigned: October 6th, 2017 Due: October 20th, 2017

Getting Started: In this assignment, you will train and evaluate different classification models on one dataset. **Please install Python 3.6 via Anaconda on your personal machine.** Download the homework file HW02.zip via Moodle. Unzipping this folder will create the directory structure shown below,

```
HW02
--- HW02.pdf
--- Data
    |--data_train.npz
    |--labels_train.npz
    |--data_test.npz
--- Submission
    |--Code
    |--Figures
    |--Predictions
```

The data files are in 'Data' directory respectively. You will write your code under the Submission/Code directory. Make sure to put the deliverables (explained below) into the respective directories.

Deliverables: This assignment has three types of deliverables:

- **Report:** The solution report will give your answers to the homework questions (listed below). Try to keep the maximum length of the report to 5 pages in 11 point font, including all figures and tables. Reports longer than five pages will only be graded up until the first five pages. If you have answered extra credit questions then you can exceed the five page limit but make sure put you solutions to extra credit questions on pages six and after. You can use any software to create your report, but your report must be submitted in PDF format.
- **Code:** The second deliverable is the code that you wrote to answer the questions, which will involve implementing classification models. Your code must be Python 3.6 (no iPython notebooks or other formats). You may create any additional source files to structure your code. However, you should aim to write your code so that it is possible to reproduce all of your experimental results exactly by running *python run_me.py* file from the Submissions/Code directory.

Kaggle Submissions: We will use Kaggle, a machine learning competition service, to evaluate the performance of your classification models. You will need to register on Kaggle using a *.edu email address to submit to Kaggle. Please use an appropriate user name for us to map Kaggle submissions to Moodle. You will generate test prediction files, save them in Kaggle format (helper code provided called Code/kaggle.py) and upload them to Kaggle for scoring. Your scores will be shown on the Kaggle leaderboard. The Kaggle link for this HW is <http://www.kaggle.com/c/hw2-image-classification>.

Submitting Deliverables: When you complete the assignment, you will upload your report and your code using the Gradescope.com service. Place your final code in Submission/Code. If you generated any figures place them under Submission/Figures. Finally, create a zip file of your submission directory, Submission.zip (NO rar, tar or other formats). Upload this single zip file on Gradescope as your solution to the ‘HW02-Classification-Programming’ assignment. Gradescope will run checks to determine if your submission contains the required files in the correct locations. Finally, upload your pdf report to the ‘HW02-Classification-Report’ assignment. When you upload your report please make sure to select the correct pages for each question respectively. Failure to select the correct pages will result in point deductions. The submission time for your assignment is considered to be the later of the submission timestamps of your code and report.

Academic Honesty Statement: Copying solutions from external sources (books, internet, etc.) or other students is considered cheating. Sharing your solutions with other students is also considered cheating. Posting your code to public repositories like GitHub, stackoverflow is also considered cheating. Any detected cheating will result in a grade of -100% on the assignment for all students involved, and potentially a grade of F in the course.

Task: You are given a set of gray scale images (32×32 pixels) with one (and only one) of the following objects: horse, truck, frog, ship (labels 0, 1, 2 and 3, respectively). The goal is to train a model to recognize which of the objects is present in an image. Some samples of the training images are:




Figure 1: horse




Figure 2: truck




Figure 3: frog




Figure 4: ship

You will train different models (Decision Trees, Nearest Neighbors, Linear models, Neural Networks), compare their performances and training time, and perform model selection using cross-validation.

The dataset is already partitioned into train and test blocks. Each partition is saved in a numpy binary format (.npy) file.

- Size of training set: $20,000 \times 1024$
- Size of testing set: $3,346 \times 1024$

Loading the files (helper code provided in ‘run_me.py’) will result in a $N \times p$ matrix, where N is the number of training/testing examples respectively and $p = 1024$ is the number of features. This loads a numpy array containing values between 0 and 256 (which are normalized, ie. divide each value by 256). The labels are not provided for the images in the test set. You need to predict the test labels, kagglize your predictions, upload to Kaggle to an accuracy score. For this assignment you are allowed to use `sklearn.model_selection.*` functions.

Questions:

1. (25 points) Decision trees:

- (10) a. A labeled data set D with N samples, each of which consists of F binary features, is given. Suppose that feature f_i was chosen in the root node while building a decision tree, splitting the data in two subsets, D_0 and D_1 . Give an expression for the information gain in this case. More generally, what criterion can be used to choose which feature to use?
- (15) b. Train 5 different decision trees using the following maximum depths $\{3, 6, 9, 12, 14\}$. Using 5-fold cross-validation, estimate the out of sample error for each model, and report them using a table. How does the maximum depth of the tree affect the estimated accuracy? Explain in at most 4 sentences. Choose the model with lowest estimated out of sample error, train it with the full training set, and predict the labels for the images in the test set. Upload your predictions to Kaggle and report the accuracy on the public leaderboard. Is the predicted out of sample error close to the real one (validation set)? Make sure that your report clearly states which model was chosen and what was the predicted out of sample error for it.

2. (25 points) Nearest neighbors:

- (10) a. A labeled dataset D with N samples, each of which consists of F features, is given. Suppose that a new sample X wants to be classified using KNN, what is the time complexity of this operation if a brute force approach is used?
- (15) b. Train 5 different nearest neighbors classifiers using the following number of neighbors $\{3, 5, 7, 9, 11\}$. Using 5-fold cross-validation, estimate the out of sample error for each model, and report them using a table. Choose the model with lowest estimated out of sample error, train it with the full training set, and predict the labels for the images in the test set. Upload your predictions to Kaggle and report the accuracy on the public leaderboard. Make sure that your report clearly states which model was chosen and what was the predicted out of sample error for it.

3. (10 points) Linear model:

- (10) a. Train a linear model using $L2$ regularization, with the following regularization constants $\alpha = \{10^{-6}, 10^{-4}, 10^{-2}, 1, 10\}$, and with hinge and logistic regression loss (you will train 10 classifiers). Look at `sklearn.linear_model.SGDClassifier` to answer this question. Using 5-fold cross-validation, estimate the out of sample error for each model, and report them using a table. Choose the model with lowest estimated out of sample error, train it with the full training set, and predict the labels for the images in the test set. Upload your predictions to Kaggle and report the accuracy on the public leaderboard. Make sure that your report clearly states which model was chosen and what was the predicted out of sample error for it.

4. (40 points) Neural Network: You will train several neural networks (NN) with one hidden layer, as the one shown in Figure 5¹.

¹Image extracted from <http://neuralnetworksanddeeplearning.com/chap5.html>




Figure 5: Sample Neural Network

- The size of the input layer, I , is determined by the number of features, $\#I = 32 \times 32 = 1024$
- The size of the hidden layer, h , will be set to M , for $M \in \{5, 40, 70\}$
- The size of the output layer, O , is determined by the number of classes. For this dataset this is set to 4

The input layer is densely connected to the hidden layer. The hidden layer is densely connected to the output layer. Given an input x the output is given by,

$$\vec{f}(x) = \mathbf{c} + V h_x \quad (1)$$

$$h_x = \sigma(b + Wx) \quad (2)$$

$$\sigma(s) = \tanh(s) \quad (3)$$

$$\frac{d\sigma(s)}{ds} = 1 - \tanh(s)^2 \quad (4)$$

Where \mathbf{b} is the bias in the hidden layer, W the matrix that contains the weights of the connections between the input and the hidden layer, V contains the weights between the hidden and output layers and c is the bias in the output layer. Note that $\vec{f}(x)$ is a vector of length matching the number of outputs. Finally, the loss function is defined as

$$L(x, y) = -\vec{f}(x)_y + \log \sum_{j=0}^3 \exp \vec{f}(x)_j \quad (5)$$

where y is the correct label associated to the input x and $\vec{f}(x)_i$ the output of the i -th neuron.

For this HW assignment we would like you to compute the partial derivatives by hand per example to better understand the inner workings of error backpropagation algorithm. But computing the partial derivatives per

example and updating the weights of the NN is time consuming from a programming stand point. In order to balance these two goals, we have designed the HW assignment such that you will compute the partial derivatives for a single example and compare your outputs to an automatic differentiation toolbox called *autograd*. You can install autograd by typing,

```
pip install autograd
```

in a terminal or Anaconda. More information on the toolbox is available here: <https://github.com/HIPS/autograd/blob/master/README.md>. For part (a) you will compare the outputs of your partial derivatives to that of autograd for a single data example. For part (b) you will only use the autograd toolbox to train a neural network. The autograd toolbox computes the partial derivatives for all 20,000 train examples simultaneously hence speeding up the training of your NN. Helper code is given for both parts of this question.

(20) a. Implement backprop for one data example: The expressions of the gradient of the loss function for a single data example with respect to the four parameters (c, V, b, W, h) is given in the lecture handout (in the week of October 9th to 13th). Use these partial derivatives to,

1. **Write python Code computing partial derivatives:** Write code in python implementing the four partial derivatives. Make sure the dimensions match. To get you started we have provided some helper code in function 'NN_gradient_one_sample.py'. Please write your partial derivatives in the 'partial_derivatives' function and list your code in your HW2 report as,

```
dLdc = ...  
dLdV = ...  
dLdb = ...  
dLdW = ...
```

Note that we would like you to list only the four partial derivatives. Please do not list import, print, function definitions, comments or other irrelevant pieces of code. Ideally the code listing will be between four and ten lines.

2. **Output values of partial derivatives for one data example:** Copy and paste the output of 'NN_gradient_one_sample.py' in your HW2 report. Sample output should look like,

```
Loss = 123...  
dLdc, Autograd  
456...  
dLdc, partial derivative  
456...  
etc,
```

(20) b. Train NN using autograd toolbox: Train three neural networks using $\{5, 40, 70\}$ as values for M . Couple of points on the specifics of the NN to be used for this question,

1. Use 1000 epochs (an epoch is defined as one forward and one backward pass on all training examples)
2. Fixed learning rate of 0.0001
3. Use 'tanh' functions for the hidden units

4. Use fixed momentum of 0.1
5. Regularization of the logistic loss function with a constant fixed penalty of 10

To get you started we have provided a helper file ‘NN_one_layer.py’ with these default settings. In this file, we have also included autograd functions to compute the gradients for the entire train dataset. Add your code to this file to,

1. Plot the mean training logistic loss (Equation 5) as a function of the number of epochs for each NN (one figure with 3 lines, make sure to label the axes and include legends; look at Code/plotting.py to get started)
2. Report, using a table, the training time (in milliseconds) for each NN
3. Using just one train-validation stratified split (uniform number of examples from each class in both train and validation sets) of 0.8/0.2, estimate the validation set error for each NN, and report them using a table. How does the number of hidden units affect the estimated accuracy? Explain in at most 3 sentences. Choose the NN with lowest estimated validation error, train it with the full training set, and predict the labels for the images in the test set. Upload your predictions to Kaggle and report the accuracy on the public leaderboard. Is the predicted validation error close to the real one? Make sure that your report clearly states which model was chosen and what was the predicted validation error for it.

Extra Credit: Finally, here are some extra-credit problems. These are far more difficult than the above problems and have very small point values. These are also deliberately more open-ended, leaving you more space for creativity. As a result, you will need to carefully describe exactly what you did for each problem. To maximize your score with limited time, you should make sure the above problems are done thoroughly and ignore these. We will be very stingy in giving credit for these problems– do them only for the glory, and only at your own risk!

5. (5 points) Extra credit: For the previous question you were asked to train a neural network with only one hidden layer. Derive the expression of the gradient with respect to the parameters for a Neural Network with 3 hidden layers, and train and evaluate the system using the training and testing dataset. Report the results obtained (accuracy and training time) using tables.

6. (5 points) Extra credit: Visualize the weights in the input and hidden layer and write your interpretation of the weight matrices. This is particularly interesting in image datasets to see the correspondence between the input images and what the NN is learning.